RESEARCH REPORT — APRIL 2026

# Speed as an AI Advantage

Prepared in connection with a private GILD executive dinner. Presented in partnership with Tecla helping AI-forward companies hire elite talent.

**88%**
USE AI

**46%**
REACH PRODUCTION

**6%**
SEE IMPACT

# Table of Contents


01 A Research-Backed, Operator-Grade Asset

02 Speed as an AI Advantage: Three Critical Observations

03 The Execution Divide

04 From AI as a Project to AI as a Continuous System

05 The Speed Architecture: Four Dimensions of Execution Advantage

06 What Is Working in 2026: Implementation Paths

07 Most AI Speed Failures Are Organizational, Not Technical

08 How Speed Compounds Across the Organization

09 Leadership Checklist: Assess Your Execution Posture

10 Five Experiments to Build Execution Velocity

SECTION 01

# A Research-Backed, Operator-Grade Asset

This research report is published after each GILD event as an additional value-add for attendees and registrants. It is designed to help CEOs, founders, and senior operators move faster inside their organizations on the topic discussed.

- A research-backed, operator-grade asset leaders can confidently share internally
- Insights synthesized from external research and translated into practical guidance
- A standalone document that stands on its own merit

> For CEOs and founders, the implication is direct: the question is no longer whether to accelerate. It is how to build the organizational infrastructure that makes acceleration sustainable.

SECTION 02

# Speed as an AI Advantage: Three Critical Observations

AI strategy in 2026 is defined less by model access and more by the organizational systems that determine how quickly insight becomes action and how consistently action generates learning.

---

> The execution divide in 2026 is not about access to AI, it is about the distance between pilot and production.

*Source synthesis: McKinsey State of AI 2026, IDC CIO Playbook 2026, PwC 29th Annual Global CEO Survey*

SECTION 03

# The Execution Divide

Pilot vs. Production Gap — The gap is not access. It is execution.

**88%**

**USE AI**

of organizations use AI in at least one function — access is nearly universal

**46%**

**REACH PRODUCTION**

of pilots ever reach production — fewer than half make it past the pilot stage

**6%**

**SEE IMPACT**

report meaningful enterprise-wide financial impact — most investment is not compounding

| WHERE MOST ORGANIZATIONS ARE | WHERE COMPOUNDERS ARE |
| --- | --- |
| AI in at least one function | AI in production across core workflows |
| 46% of pilots reach production | Pilots ship in weeks, not quarters |
| Speed measured by announcements | Speed measured by learning cycle time |
| Governance blocks experimentation | Governance enables pre-authorized experiments |
| Hiring for AI tool familiarity | Hiring for critical thinking and orchestration |

SECTION 04

# From AI as a Project to AI as a Continuous System

The traditional operating model treated AI deployment as a project with defined phases. The new model treats AI deployment as a continuous system.

| TRADITIONAL OPERATING MODEL | AI-NATIVE EXECUTION MODEL |
|---|---|
| AI as a tooling upgrade added to existing workflows | AI embedded into redesigned workflows and decision systems |
| Speed measured by shipping features | Speed measured by time from hypothesis to validated learning |
| Bottleneck: gaining internal approval to move | Bottleneck: absorbing and validating change safely |
| Talent hired for AI familiarity | Talent hired for critical thinking + AI orchestration |
| Pilots evaluated quarterly by committees | Experiments evaluated weekly by defined owners and metrics |
| Advantage measured by AI spend | Advantage measured by learning loop velocity |

> In the old model, speed was a value. In the new model, **speed is an architecture.**

SECTION 05

# The Speed Architecture

Four Dimensions of Execution Advantage — AI-driven execution speed is the output of four interdependent organizational dimensions.

01

### Decision Rights

Clear ownership for fast approvals. Pre-authorized categories of AI experiments that teams can launch without case-by-case approval.

02

### Talent & Judgment

Hire for critical thinking first. 73% of talent leaders rank problem-solving as the #1 skill needed; AI technical skills rank fifth.

03

### Measurement & Feedback

Metrics and loops to iterate quickly. Outcomes measured within 48–72 hours of deployment, not at the next quarterly review.

04

### Platform & Integration

Reusable infrastructure for rapid deployment. High performers are nearly three times as likely to have redesigned their workflows.

> **Key Insight:** 73% of talent leaders rank critical thinking and problem-solving as the #1 skill needed; AI technical skills rank fifth. Hire for judgment, not just tool familiarity.

SECTION 06

# What Is Working in 2026: Implementation Paths

Six best practices from high-velocity teams

01

### Define the experiment before the tool

High-velocity teams specify what they are trying to learn, who owns the decision, what metric determines success, and what the review cadence is, before any AI capability is selected or deployed.

02

### Separate permission to experiment from permission to scale

Maintain a sharp distinction between low-cost experiments and scaled deployments. Conflating the two slows both.

03

### Measure learning speed, not just shipping speed

The most useful metric is not how many features shipped, it is how quickly the organization moved from hypothesis to validated learning.

04

### Build talent for orchestration, not just operation

Most people can learn to use AI tools, but far fewer can rigorously evaluate their output. Hiring strategies that prioritize critical thinking over tool certification produce teams that compound.

05

### Treat hesitation as a cost center

Organizations that hesitated on AI talent are now paying 15–20% premiums for equivalent skills. Delay has a price, and in 2026, that price is visible across talent, capability building, and workflow redesign.

06

### Redesign workflows before layering AI

High performers are nearly three times as likely to have fundamentally redesigned their workflows rather than simply layered AI onto existing ones. Deployment without redesign produces no compounding.

SECTION 07

# Most AI Speed Failures Are Organizational, Not Technical

Understanding these failure modes is the first step to avoiding them.

| FAILURE MODE | MITIGATION |
| --- | --- |
| Mistaking motion for velocity, shipping fast without learning fast | Define success criteria and measurement cadence before each experiment begins |
| Governance designed to slow movement rather than enable safe movement | Redesign governance as pre-authorization of categories, not case-by-case approval |
| Hiring for AI tool familiarity instead of critical thinking and orchestration capability | Audit hiring criteria against the actual skills needed to evaluate and direct AI outputs |
| Compressing decision timelines without compressing information flows | Invest in feedback infrastructure that surfaces signals faster, not just faster meetings |
| Building speed in one function while the rest of the organization operates at legacy pace | Sequence AI deployment to create cross-functional learning density, not isolated pockets |
| Treating AI talent as a commodity hire rather than a strategic architecture decision | Move on talent faster; demand outstrips supply by 3.2 to 1, and waiting increases cost |

SECTION 08

# How Speed Compounds Across the Organization

AI execution advantage accumulates across five organizational layers. Each layer feeds the next.

**5 Compounding Advantage**
Market position that widens with every completed cycle

**4 Deployment Velocity**
Production-grade AI embedded in core workflows

**3 Learning Loop Design**
Short cycles from action to insight to next decision

**2 Talent Architecture**
Critical thinkers who orchestrate AI and evaluate its outputs

**1 Infrastructure & Governance**
Decision rights, experiment authorization, fast feedback tools

FOUNDATION LAYER

> **The advantage lives in the loops, not in the technology.** Each completed cycle makes the next cycle faster and more accurate. This is the layer competitors cannot replicate by buying the same tools.

SECTION 09

# Leadership Checklist: Assess Your Execution Posture

### Decision Velocity

- [ ] Pre-authorized categories of AI experiments that teams can launch without case-by-case approval
- [ ] Decision rights for AI initiatives clearly defined by scope and risk level
- [ ] Average time from idea to authorized experiment is known and tracked
- [ ] Slowest approval processes are visible to leadership and actively being reduced

### Talent Architecture

- [ ] Hiring criteria for AI-related roles prioritize critical thinking and orchestration over tool certifications
- [ ] Moving fast on AI talent, the cost of delay in this market is understood
- [ ] Team members responsible for AI initiatives can evaluate AI outputs, not just produce them
- [ ] Identified which AI capabilities require permanent hires vs. flexible staffing models

### Learning Loop Design

- [ ] Every active AI experiment has pre-defined success criteria established before launch
- [ ] Outcomes measured within 48–72 hours of deployment, not at the next quarterly review
- [ ] Structured weekly or bi-weekly cadence for reviewing experiment results
- [ ] Feedback loops are shorter than governance cycles

### Governance as Enabler

- [ ] Governance is designed to pre-authorize safe experiments, not audit after the fact
- [ ] Individual teams can act without seeking approval for each decision within defined parameters
- [ ] Risk thresholds are clearly defined and communicated across the organization

SECTION 10

# Five Experiments to Build Execution Velocity

01

### Decision Speed Audit

**Owner:** CEO / COO
**Metric:** Average days from idea to authorized experiment
Map the actual time from identified AI opportunity to authorized experiment across three recent initiatives. Identify the primary bottleneck and redesign that specific step.

02

### Learning Loop Compression Pilot

**Owner:** VP of Product or Engineering
**Metric:** Time from deployment to team decision on next action
Choose one active AI initiative and compress its feedback cycle. Pre-define success criteria before launch. Measure within 72 hours.

03

### Talent Architecture Review

**Owner:** CTO / Head of Talent
**Metric:** Ratio of orchestration-capable hires to tool-specialist hires
Audit current and planned AI hires against a skills framework that prioritizes critical thinking, output evaluation, and workflow orchestration.

04

### Governance Redesign Sprint

**Owner:** General Counsel / CTO
**Metric:** Number of experiments launched without individual approval
Identify the three most common approval steps that slow AI experiments. For each, design a pre-authorization category that removes the approval step for experiments below a defined risk and cost threshold.

05

### Speed-to-Learning Benchmarking

**Owner:** CEO / CTO
**Metric:** Days from hypothesis to validated learning in pilot function
Define a baseline for your current learning loop velocity. Set a 90-day

target to reduce that time by 25% in one function.

06

### Feedback Infrastructure Build

**Owner:** COO / Head of Operations
**Metric:** Time from experiment result to organization-wide visibility
Design a lightweight system that surfaces experiment results, signals, and decisions in real time across teams. Standardize how outcomes are captured, shared, and reviewed.

CLOSING

# Closing Perspective

> "The work is unglamorous. Decision frameworks. Experiment protocols. Talent criteria. Governance redesign."

**But this is what compounding looks like before it becomes visible.**

The organizations pulling ahead are not necessarily moving faster in any given moment. They are moving *consistently*. Each experiment generates learning. Each learning generates a better next experiment. Each better experiment generates a wider gap between their capability and their competitors'.

This is not a temporary lead driven by early access to a technology. It is an **architectural lead**, built into the decision systems, the talent architecture, the feedback infrastructure, and the governance design of the organization.

**The window to build this architecture is narrowing.**


© 2026 GILD. All rights reserved.